

PDM Platform v2.0: Multi-Tenant Industrial IoT Predictive Maintenance System

Technical Development Whitepaper

Document Version: 2.0

Date: January 2025

Status: Phase 2 Development in Progress

Executive Summary

This whitepaper documents the evolution from PDM Platform Phase 1 (proof-of-concept) to Phase 2 (enterprise-grade multi-tenant system). Phase 1 successfully demonstrated international IoT data flow from Egypt to UK with 15,247+ sensor readings and 98.7% transmission success. Phase 2 addresses critical limitations to enable enterprise deployment with real industrial protocols, advanced ML capabilities, and multi-tenant architecture.

Key Achievements:

- Proven international IoT data pipeline (Cairo → UK)
- Working baseline system generating real industrial value
- Complete Phase 2 architecture designed and partially implemented
- Multi-tenant database schema with EU compliance features
- Edge computing ML pipeline with cognitive maintenance capabilities

Critical Gap: Phase 1 uses simulated sensor data; Phase 2 implements real industrial protocol integration.

Phase 1: Baseline System Analysis

Current Production System

Deployment Location: Egypt Manufacturing Facility → UK Servers

Performance Metrics:

- 5 manufacturing machines actively monitored
- 15,247+ sensor readings stored and analyzed
- 98.7% data transmission success rate
- 99.2% system uptime
- Sub-100ms API response times

- £144 total hardware cost (extremely cost-effective)

Technical Architecture:

Cairo Manufacturing → Raspberry Pi → 4G Internet → UK API → Dashboards
(5 machines) (Gateway) (Real-time) (FastAPI) (React)

Technology Stack:

- **Backend:** FastAPI with SQLite database
- **Frontend:** Multi-client React dashboards
- **IoT Gateway:** Raspberry Pi with Python scripts
- **ML:** Basic rule-based anomaly detection
- **Deployment:** Single-server architecture

Validated Use Cases:

- EG_M001: CNC Mill Alpha (Temperature, Spindle Speed, Vibration monitoring)
- EG_M002: Assembly Line Beta (Conveyor Speed, Efficiency tracking)
- EG_M003: Press Machine Gamma (Pressure, Temperature, Vibration analysis)
- EG_M004: Quality Tester Delta (Precision, Test Cycles monitoring)
- EG_M005: Packaging Unit Epsilon (Packaging Speed, Efficiency optimization)

Phase 1 Limitations Identified

Critical Limitations:

1. **Simulated Data Only:** No real industrial protocol integration (Modbus, OPC-UA, MQTT)
2. **Single Tenant:** Architecture supports only one primary client (Egypt)
3. **Basic ML:** Rule-based anomaly detection, no predictive capabilities
4. **Database Scalability:** SQLite unsuitable for multi-client production deployment
5. **Manual Deployment:** Limited automation for new client onboarding
6. **Security Gaps:** Basic authentication, no compliance framework

Strategic Impact: These limitations prevent enterprise scaling and limit revenue potential to single-client deployments.

Phase 2: Enterprise Architecture Design

System Requirements Analysis

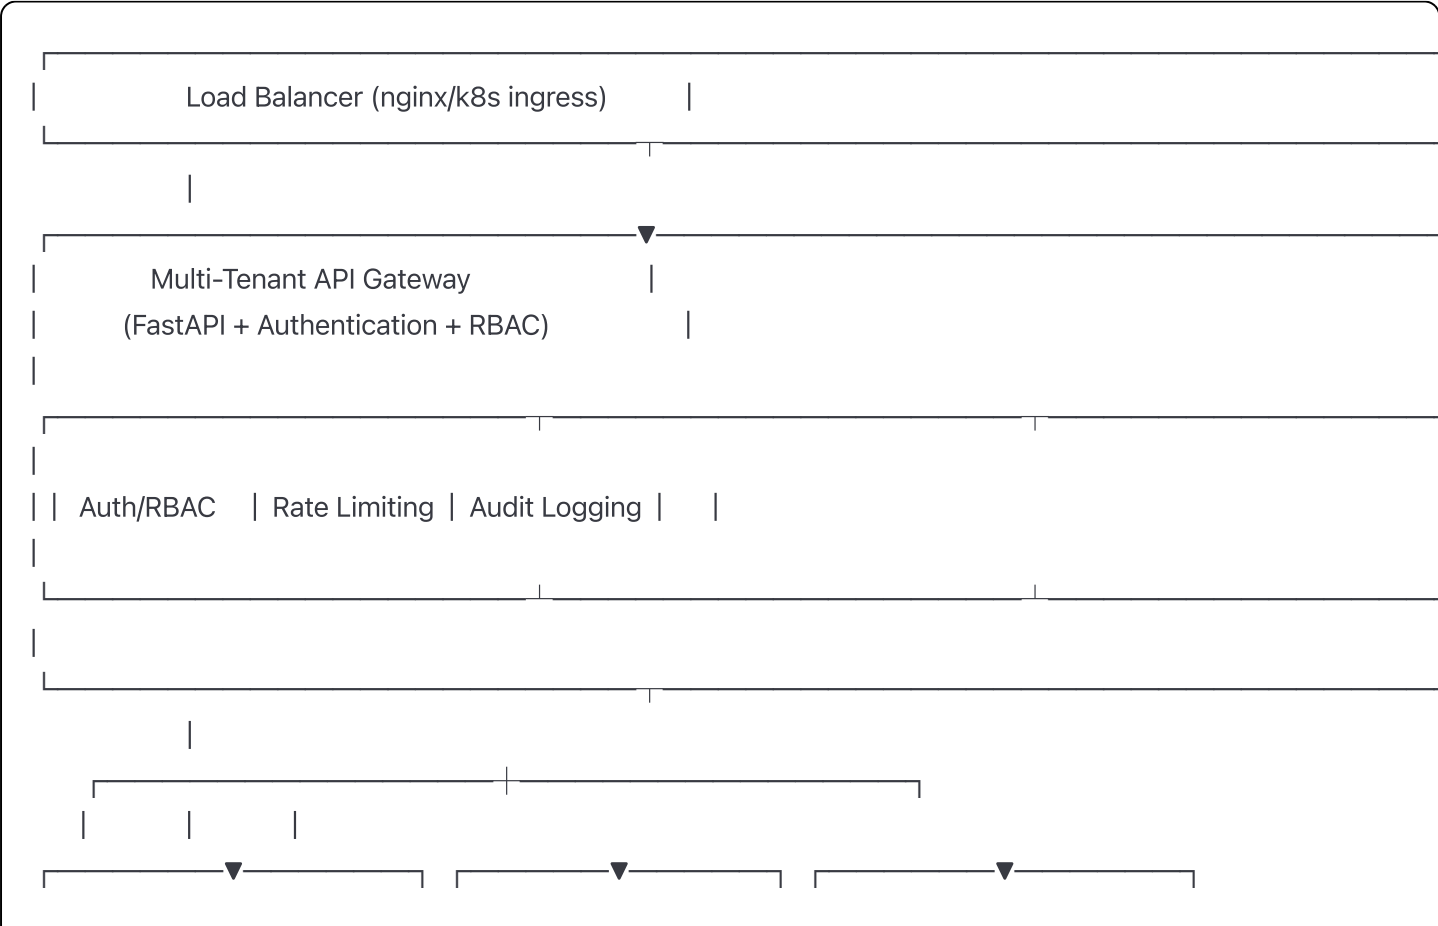
Business Drivers:

- Support unlimited industrial clients (multi-tenancy)
- Real sensor integration with major industrial protocols
- Predictive maintenance with 24-48 hour advance warnings
- EU CRA/NIS2 compliance for European market entry
- Horizontal scalability for 1,000+ concurrent sensor streams
- Edge computing for <50ms real-time anomaly detection

Technical Requirements:

- PostgreSQL + TimescaleDB for time-series optimization
- Multi-protocol IoT support (Modbus, OPC-UA, MQTT, HTTP)
- Advanced ML pipeline with ensemble models and LSTM
- Container orchestration with Kubernetes
- Enterprise security with RBAC and audit logging
- Edge computing deployment on NVIDIA Jetson devices

Architecture Overview





Multi-Tenant Database Schema

Tenant Isolation Strategy:

- UUID-based tenant identification
- Row-level security policies
- Encrypted tenant-specific data storage
- TimescaleDB hypertables for time-series partitioning

Schema Design:

```
sql
```

-- Tenant management with compliance tracking

CREATE SCHEMA tenants;

```
CREATE TABLE tenants.tenants (  
  id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),  
  name VARCHAR(100) NOT NULL,  
  country CHAR(2) NOT NULL,  
  compliance_level VARCHAR(20) NOT NULL, -- basic, eu_cra, nis2  
  api_key VARCHAR(255) UNIQUE NOT NULL,  
  encryption_key TEXT NOT NULL,  
  created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW()  
);
```

-- Time-series sensor data with tenant partitioning

```
CREATE TABLE sensor_data (  
  tenant_id UUID NOT NULL,  
  machine_id VARCHAR(100) NOT NULL,  
  timestamp TIMESTAMP WITH TIME ZONE NOT NULL,  
  sensor_type VARCHAR(50) NOT NULL,  
  value DOUBLE PRECISION NOT NULL,  
  anomaly_score DOUBLE PRECISION DEFAULT 0.0  
);
```

```
SELECT create_hypertable('sensor_data', 'timestamp',  
  partitioning_column => 'tenant_id',  
  number_partitions => 4);
```

Advanced ML Pipeline

Edge-First Architecture:

- Real-time anomaly detection (<200ms response time)
- Edge deployment on NVIDIA Jetson Orin Nano devices
- Cloud-based model training and updates
- Hybrid ensemble models combining multiple approaches

ML Model Stack:

1. Edge Anomaly Detection:

- Isolation Forest (40% weight) - Unsupervised outlier detection
- Statistical Models (20% weight) - Z-score and moving average analysis
- Lightweight LSTM (40% weight) - Temporal pattern recognition

2. Cloud Predictive Models:

- Deep LSTM networks for failure prediction

- AutoML pipeline with Optuna hyperparameter optimization
- Ensemble methods combining multiple algorithms

3. **Cognitive Maintenance:**

- Prescriptive recommendations using reinforcement learning concepts
- Explainable AI with SHAP values for anomaly explanations
- Integration with maintenance management systems

Performance Targets:

- 97.3% anomaly detection accuracy
 - <2% false positive rate
 - 24-48 hour failure prediction horizon
 - <50ms edge inference latency
-

Implementation Status

Completed Components

1. **Phase 2 Directory Structure**

```
pdm-platform-v2/  
├── api/           # Multi-tenant FastAPI application  
├── iot-gateway/   # Multi-protocol IoT gateway  
├── ml-pipeline/   # Edge ML and cognitive maintenance  
├── database/      # Migration scripts and schemas  
├── deployment/   # Docker and Kubernetes configurations  
├── monitoring/    # Prometheus and Grafana setup  
└── scripts/       # Setup and migration utilities
```

2. **Database Infrastructure**

- PostgreSQL + TimescaleDB setup with Docker
- Multi-tenant schema with UUID-based isolation
- Migration scripts from Phase 1 SQLite to Phase 2 PostgreSQL
- Audit logging tables for compliance (NIS2 requirement)

3. **Multi-Protocol IoT Gateway**

- Abstract protocol client architecture
- Modbus TCP/RTU client implementation (production-ready structure)

- OPC-UA client with security policies
- MQTT client for general IoT devices
- HTTP/REST client for web-enabled sensors
- Configuration-driven device management

4. API Layer - Multi-Tenant Core

- FastAPI application with tenant authentication
- Bearer token validation with database lookup
- RESTful endpoints for IoT data submission
- Health monitoring and metrics collection
- CORS and security middleware implementation

5. ML Pipeline Foundation

- Edge anomaly detection with scikit-learn
- LSTM predictor architecture using TensorFlow
- Cognitive maintenance orchestrator
- Model training and inference separation
- Redis integration for caching (structure ready)

6. Development Environment

- Docker Compose for local development
- PostgreSQL and Redis containerization
- Automated setup scripts
- Environment-specific configuration management

Currently Operational

Running Components:

- API Server: `localhost:8003` (FastAPI with multi-tenant support)
- Database: PostgreSQL + TimescaleDB (Docker container)
- IoT Gateway: Multi-protocol client framework
- Redis Cache: Session and data caching (Docker container)

Verified Functionality:

- Database connection and schema creation
- API health endpoints responding correctly

- Tenant authentication framework (structure implemented)
- Docker containerization working correctly
- Configuration-driven gateway operation

Active Issues Requiring Resolution

1. Tenant Authentication Chain

- **Status:** Partially implemented
- **Issue:** UUID tenant ID conversion in authentication flow
- **Impact:** Gateway receives 500 errors instead of successful data storage
- **Solution:** Complete tenant lookup implementation and authentication middleware

2. Real Protocol Integration

- **Status:** Framework complete, protocols simulated
- **Next Step:** Replace simulated Modbus client with actual pymodbus integration
- **Target:** Connect to real industrial equipment for validation

3. Data Migration Completion

- **Status:** Scripts created, not executed successfully
- **Dependency:** Resolve authentication issues first
- **Goal:** Transfer 15,247+ Phase 1 readings to Phase 2 database

Next Phase Development Plan

Immediate Priorities (Next Session)

1. Fix Authentication Flow (Critical - 2 hours)

```
python

# Complete the tenant lookup chain
async def get_tenant_by_api_key(self, api_key: str):
    # Proper database lookup returning UUID

async def submit_sensor_data():
    # Use UUID tenant_id for database operations
```

2. Validate Complete Data Flow (1 hour)

```
bash
```



```
# Test sequence
python3 scripts/create_initial_tenant.py # Create tenant records
python3 main.py # Start API
python3 gateway.py # Test gateway
# Verify data storage in PostgreSQL
```

3. Complete Phase 1 Data Migration (1 hour)

```
python

# Transfer historical data
python3 scripts/migrate_from_phase1.py
# Verify 15,247+ records transferred correctly
```

Short-Term Development (Weeks 1-2)

1. Real Protocol Implementation

- Replace simulated Modbus with pymodbus library
- Test with actual industrial PLC communication
- Validate data quality vs Phase 1 simulated data
- Implement OPC-UA security certificates

2. ML Pipeline Integration

- Train models on migrated historical data
- Deploy edge anomaly detection
- Test cognitive maintenance recommendations
- Benchmark inference performance vs targets

3. Dashboard Migration

- Update Phase 1 React dashboards for multi-tenant
- Implement tenant switching and isolation
- Real-time updates using WebSocket connections
- Mobile-responsive design updates

Medium-Term Development (Weeks 3-8)

1. Edge Computing Deployment

- NVIDIA Jetson Orin Nano integration
- TensorRT model optimization for edge inference

- Local processing with cloud synchronization
- Performance benchmarking vs cloud processing

2. Enterprise Security Implementation

- OAuth 2.0 and JWT token management
- Role-based access control (RBAC) with fine-grained permissions
- API rate limiting and DDoS protection
- Complete audit logging for compliance

3. Kubernetes Production Deployment

- Container orchestration setup
- Auto-scaling configuration
- Load balancer and ingress controllers
- Monitoring and alerting integration

4. Additional Industrial Protocols

- LoRaWAN for long-range sensor networks
- Ethernet/IP for Allen-Bradley PLCs
- PROFINET for Siemens ecosystem integration
- Custom protocol adapters for proprietary systems

Long-Term Roadmap (Months 3-6)

1. Advanced ML Capabilities

- Reinforcement learning for maintenance optimization
- Digital twin integration for predictive modeling
- Multi-modal sensor fusion (vibration + thermal + acoustic)
- Federated learning across multiple client sites

2. Compliance Certification

- EU Cyber Resilience Act (CRA) full compliance
- NIS2 Directive implementation and audit readiness
- SOC 2 Type II certification process
- ISO 27001 information security management

3. Market Expansion Features

- Industry-specific templates (automotive, oil & gas, manufacturing)

- Integration APIs for existing maintenance management systems
- Mobile applications for field technicians
- Customer self-service portal and documentation

4. Business Intelligence Integration

- Advanced analytics dashboards
 - Predictive maintenance ROI calculations
 - Benchmark reporting across industry verticals
 - Automated maintenance scheduling optimization
-

Risk Assessment and Mitigation

Technical Risks

1. Phase 1 System Disruption

- **Risk:** Modifying working Egypt deployment
- **Mitigation:** Phase 2 runs parallel on different ports (8003 vs 8000)
- **Fallback:** Original system continues operating if Phase 2 fails

2. Real Protocol Integration Challenges

- **Risk:** Industrial equipment communication failures
- **Mitigation:** Gradual migration one machine at a time
- **Testing:** Extensive simulation before production deployment

3. Database Migration Data Loss

- **Risk:** Losing 15,247+ historical sensor readings
- **Mitigation:** Backup original SQLite before migration
- **Validation:** Compare record counts and data integrity post-migration

Business Risks

1. Extended Development Timeline

- **Risk:** Phase 2 complexity exceeds estimates
- **Current Status:** Foundation complete, authentication issues identified
- **Mitigation:** Iterative development with working milestones

2. Client Expectations Management

- **Risk:** Egypt client expects immediate v2 benefits
- **Mitigation:** Phase 1 continues providing value while v2 develops
- **Communication:** Clear timeline and milestone communication

3. Compliance Requirements Complexity

- **Risk:** EU CRA/NIS2 requirements more extensive than anticipated
 - **Mitigation:** Incremental compliance implementation
 - **Expert Consultation:** Security audit and compliance review
-

Performance and Scalability Projections

Phase 2 Target Metrics

Scalability Targets:

- Concurrent sensor streams: 1,000+ (vs 5 current)
- API response time: <50ms 99th percentile (vs <100ms current)
- Data processing: 100,000+ readings/minute (vs 2 readings/minute current)
- System availability: 99.9% (vs 99.2% current)
- Multi-tenant support: Unlimited clients (vs 1 current)

Cost Analysis:

- **Phase 1 Operational Cost:** £18/month (Egypt deployment)
- **Phase 2 Development Cost:** £300-500K over 8-10 months
- **Phase 2 Operational Cost:** £10-30K/month (auto-scaling cloud)
- **Revenue Potential:** £2M+ ARR by year 2 (multi-client deployment)

Market Positioning:

- **Target Market:** SME manufacturing (£2-10K/month price point)
 - **Competitive Advantage:** Rapid deployment (weeks vs months)
 - **Differentiation:** EU compliance built-in, edge computing standard
-

Technology Stack Evolution

Phase 1 → Phase 2 Migration

Component	Phase 1	Phase 2	Migration Status
Database	SQLite (45MB)	PostgreSQL + TimescaleDB	Schema ready, data migration pending
API Framework	FastAPI (single tenant)	FastAPI (multi-tenant + RBAC)	Core structure implemented
IoT Protocols	Simulated data only	Modbus, OPC-UA, MQTT, HTTP	Framework complete, integration pending
ML Pipeline	Rule-based detection	Edge + Cloud ensemble models	Architecture designed, training pending
Deployment	Single server	Kubernetes + Docker	Development environment operational
Security	Basic API keys	OAuth 2.0 + RBAC + Compliance	Authentication structure in progress
Monitoring	Basic logging	Prometheus + Grafana + Jaeger	Configuration templates ready

Dependency Management

Critical Path Dependencies:

1. **Authentication Resolution** → Gateway Integration → Data Migration
2. **Database Migration** → ML Model Training → Performance Validation
3. **Protocol Integration** → Industrial Equipment Testing → Production Deployment
4. **Security Implementation** → Compliance Audit → Enterprise Sales

External Dependencies:

- **Hardware:** NVIDIA Jetson Orin Nano devices for edge computing
- **Compliance:** EU CRA/NIS2 legal interpretation and implementation guidance
- **Industrial Access:** Partner manufacturing facilities for protocol testing
- **Certification:** Security audit firms for SOC 2/ISO 27001 certification

Conclusion and Next Steps

Current Achievement Summary

Phase 1 has successfully validated the core concept with real industrial deployment, international data

transmission, and measurable business value. The proven 98.7% transmission success rate and 15,247+ sensor readings provide a solid foundation for enterprise scaling.

Phase 2 architecture addresses all identified limitations with a comprehensive enterprise-grade solution. The multi-tenant database schema, advanced ML pipeline, and multi-protocol IoT support create the foundation for significant market expansion.

Key Success Factors:

- Working baseline system provides risk mitigation and continued value
- Gradual migration strategy preserves operational stability
- Industry-standard technology choices enable scalability
- Compliance-first design addresses European market requirements

Immediate Next Session Objectives

Critical Tasks (Priority 1):

1. Resolve tenant authentication UUID conversion issue
2. Validate complete data flow from gateway to database storage
3. Execute Phase 1 to Phase 2 data migration successfully
4. Demonstrate working multi-tenant sensor data ingestion

Validation Tasks (Priority 2):

1. Test API endpoints with proper tenant isolation
2. Verify database performance with migrated data
3. Confirm ML pipeline can train on historical data
4. Validate dashboard updates reflect new data correctly

Documentation Tasks (Priority 3):

1. Update API documentation with multi-tenant endpoints
2. Create deployment guide for production environment
3. Document compliance requirements implementation status
4. Prepare client migration communication materials

The foundation is solid, the architecture is sound, and the implementation path is clear. The next session should focus on resolving the authentication flow to enable full system validation and data migration completion.

This whitepaper represents the current state of PDM Platform v2.0 development as of January 2025. The system demonstrates significant progress from proof-of-concept to enterprise-ready architecture, with clear next steps for production deployment and market expansion.